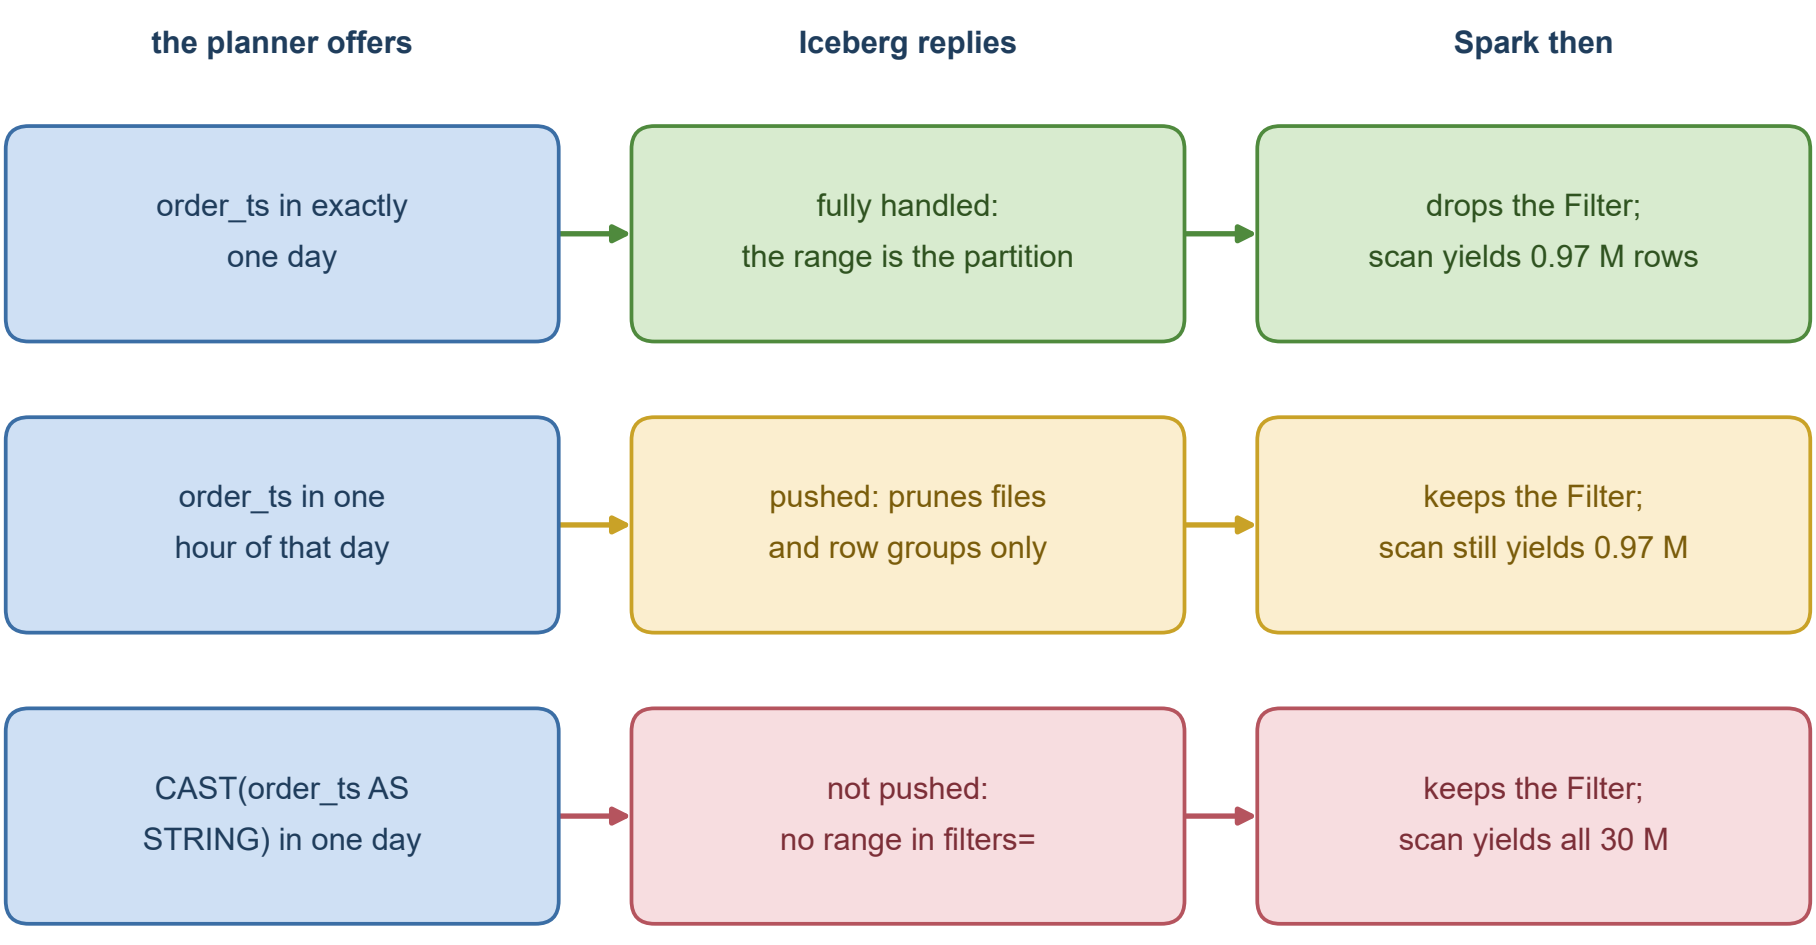
the planner offers
Iceberg replies
Spark then
order_ts in exactly one day
fully handled: the range is the partition
drops the Filter; scan yields 0.97 M rows
order_ts in one hour of that day
pushed: prunes files and row groups only
keeps the Filter; scan still yields 0.97 M
CAST(order_ts AS STRING) in one day
not pushed: no range in filters=
keeps the Filter; scan yields all 30 M