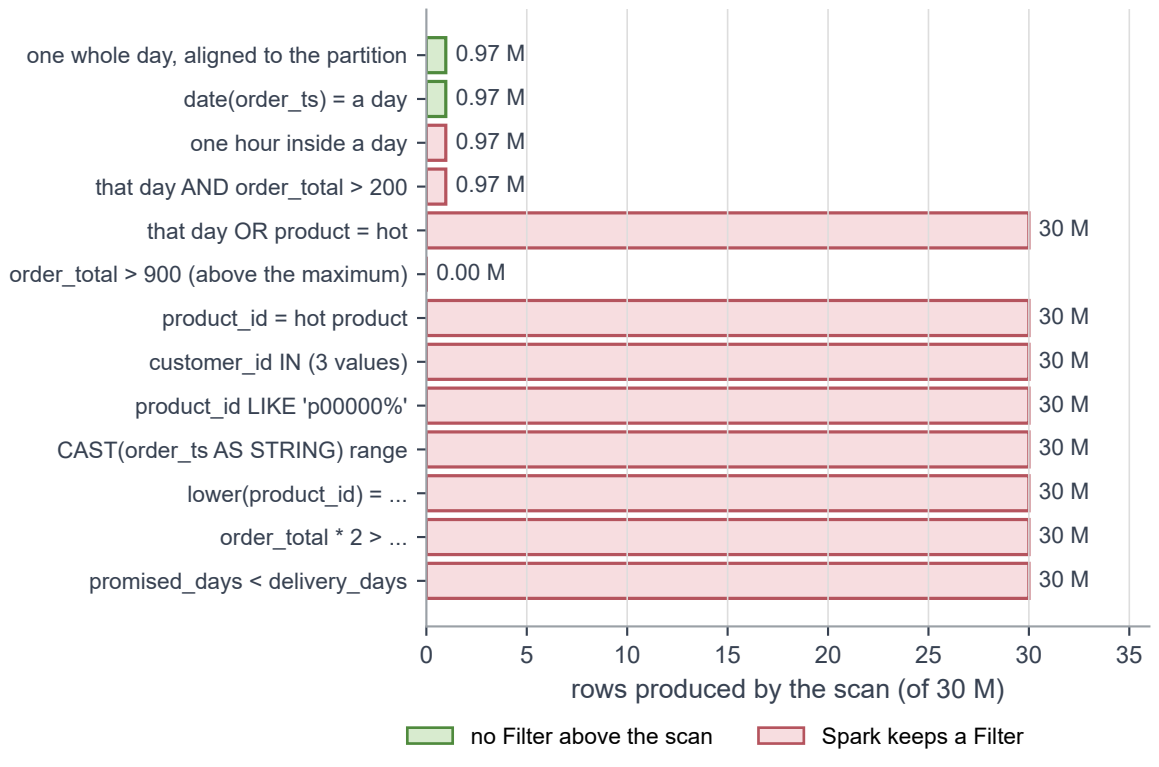
one whole day, aligned to the partition
0.97 M
date(order_ts) = a day
0.97 M
one hour inside a day
0.97 M
that day AND order_total > 200
0.97 M
that day OR product = hot
30 M
order_total > 900 (above the maximum)
0.00 M
product_id = hot product
30 M
customer_id IN (3 values)
30 M
product_id LIKE 'p00000%'
30 M
CAST(order_ts AS STRING) range
30 M
lower(product_id) = ...
30 M
order_total * 2 > ...
30 M
promised_days < delivery_days
30 M
0
5
10
15
20
25
30
35
rows produced by the scan (of 30 M)
no Filter above the scan
Spark keeps a Filter